

CS422

Database systems

Today: Data streams & AQP

Data-Intensive Applications and Systems (DIAS) Laboratory
École Polytechnique Fédérale de Lausanne

Slides adapted from presentations of the Berkeley/MIT team



Overview

Previous weeks

- Big Data infrastructures & architectural choices

This week

- Data stream processing
- Approximate query processing

Data streams management

- Traditional DBMS – data stored in **finite, persistent data sets**
- **Data Streams** – distributed, **continuous, unbounded**, rapid, time varying, noisy, ...
- **Data-Stream Management** – variety of applications
 - Real-time network analytics
 - Network security
 - Traffic engineering
 - Sensor networks
 - Financial applications
 - Telecom call-detail records
 - Web logs and clickstreams
 - Manufacturing processes
 - ...

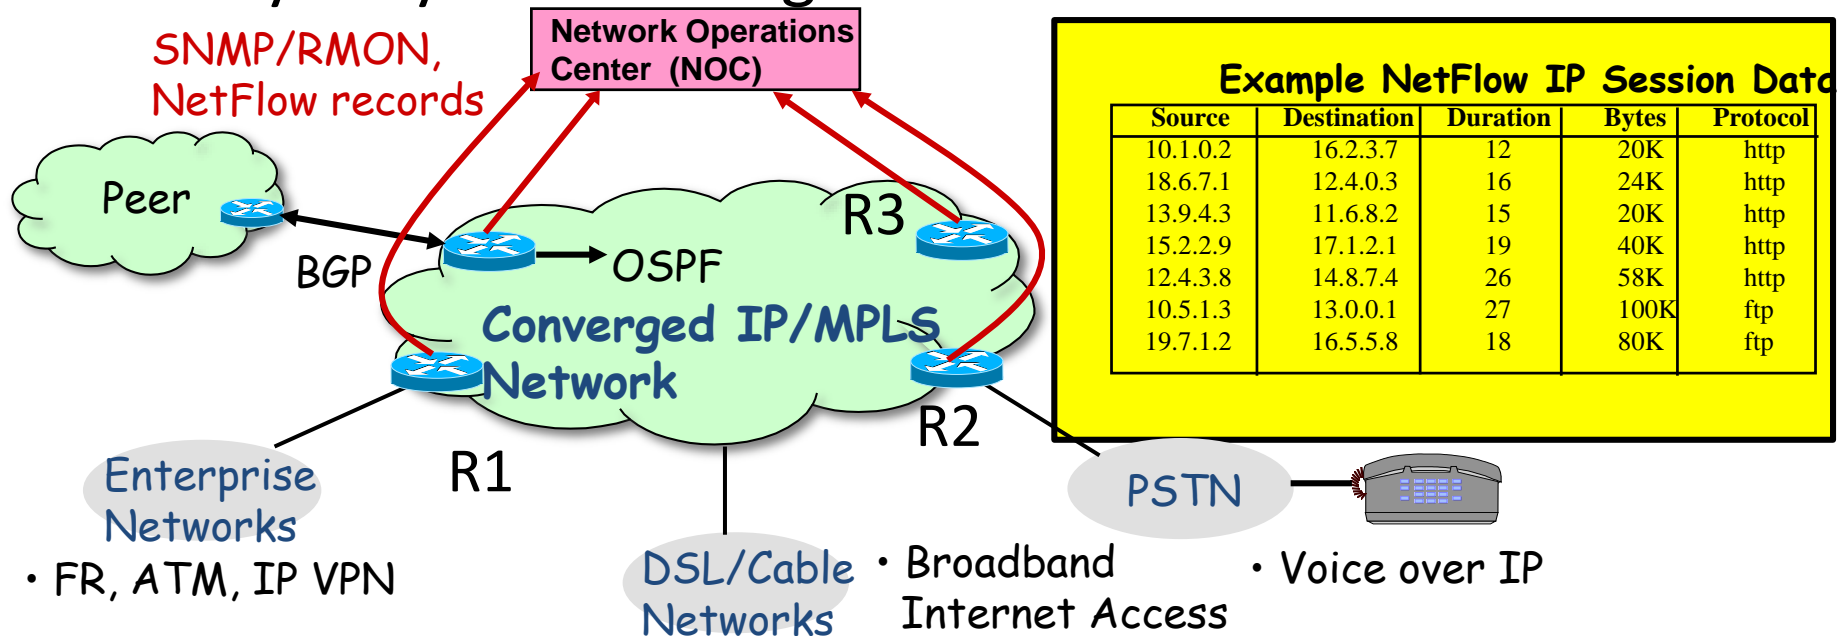
Example: stock monitoring

- Notify me when some stock price increases by at least 5% in two seconds.
- Find the top (most-frequent) 10 traded stocks in the last 10 minutes
- Notify me when correlation of two stocks over the last 10 minutes exceeds 0.6

Time (sec)	MSFT	APPL	WMT
0	65.06	121.35	66.74
1	65.06	121.36	66.75
2	65.06	121.36	66.73
3	65.07	121.36	66.72
4	65.08	121.36	66.72
5	65.07	121.35	66.71
...

Example: Real-time netw. analytics

- Which are the top (most frequent) 1000 (source, dest) ip pairs seen by R1 over the last month
- How many distinct (source,dest) pairs have been seen by R1 and R2 but not R3
- Which IP addresses receive a lot of data but send only very few messages



Example: Real-time netw. analytics (2)

Where is the challenge -- assume **packet-level** statistics

- Single 2Gb/sec link; say avg packet size is 50 bytes
- Number of packets/sec = 5 million
- If we only capture **header information** per packet:
src/dest IP, time, no. of bytes, etc. – at least 10 bytes.
 - Space per second is 50MB
 - Space per day is 4.5TB per link
 - ISPs typically have hundreds of links!
- Deep packet analysis – whole new ballgame!!

Databases Vs Data Streams

A relation is a set of tuples	A stream is a bag of tuples with partial order
Relations are persistent	Streams need to be processed in real time as tuples arrive
Interactive queries	Continuous queries
Random access to data, queries need to be processed as they arrive	Sequential access to data, random access to continuous queries
Physical database design does not change during query, queries can be unpredictable	Queries do not change, stream can be very unpredictable

Overview

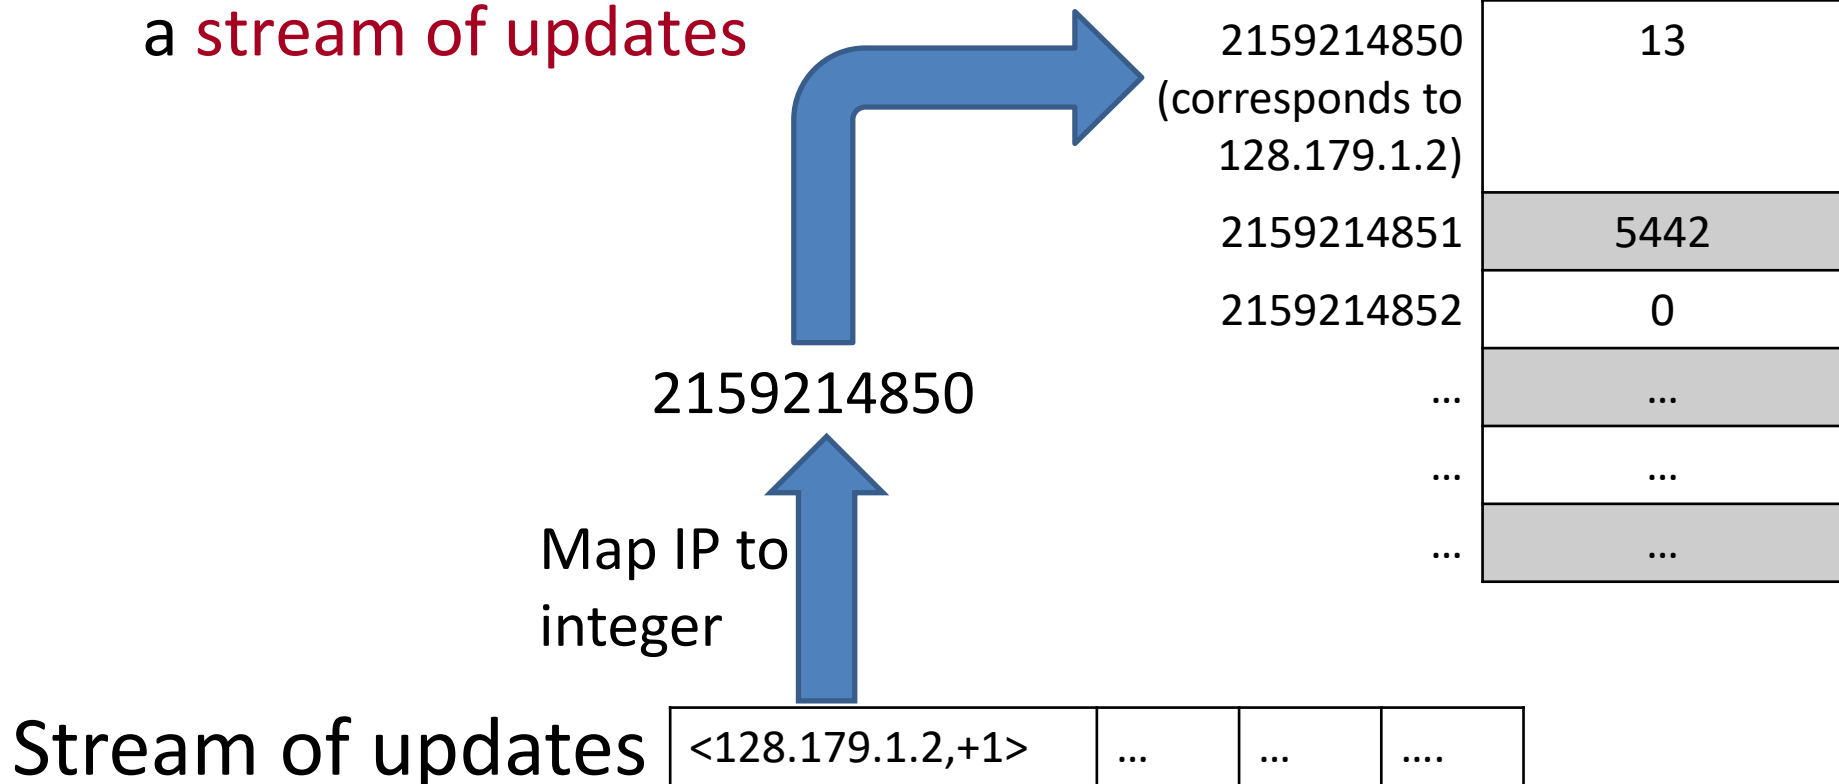
- Introduction & motivation
- **Data Stream Model**
- The data facet
- The time facet

The Data Stream Model

- **Underlying signal**: One-dimensional array $A[1...N]$ with values $A[i]$ all initially zero
 - Multi-dimensional arrays also possible
- Signal implicitly represented via a **stream of updates**
 - j -th update is $\langle k, c[j] \rangle$ implying
 - $A[k] = A[k] + c[j]$ (count $c[j]$ can be >0 or <0)
- Goal: Compute functions on $A[]$ subject to
 - Small space
 - Fast processing of updates
 - Fast function computation

The Data Stream Model – example

- **Underlying signal**: One-dimensional array $A[1...N]$
- Signal implicitly represented via a **stream of updates**



The Data Stream Model – special cases

- Cash-register model
 - $c[j]$ is always ≥ 0 (increment-only)
 - In many cases, $c[j]=1$
- Turnstile model
 - Most general streaming model
 - $c[j]$ can be positive or negative (increment or decrement)
- Time-series model
 - j -th update updates $A[j]$ (i.e., $A[j]=c[j]$)

The Data Stream Model – special cases

- Cash-register model
- Turnstile model
- Time-series model
- Difficulty varies depending on the model & problem
 - E.g., min/max in time-series Vs turnstile

The two facets of the problem

DATA

- Too much data
 - Velocity
 - Dimensions

TIME

- Interested only for parts of the stream
- Need to expire old data

Solutions of **small space/small computational complexity**

The facet of data

Stream of items:

Source	Destination	Time	Protocol	Data
10.1.0.2	16.2.3.7	1992191	http

IP network signals

- Number of active flows per source-IP address
 - 2^{32} sized array: increment & decrement (~16 GB)
- Number of packets exchanged between any two IP addresses **during the day**
 - 2^{64} sized array: increment only (64 EB!)
- Number of packets sent by each IP **in the last hour (???)**
 - 2^{32} sized array, **sliding window**

The facet of time

Queries: **Continuous**

- Defining the query range

- Landmark window




- Fixed windows



- Sliding window



- Decay model



The facet of time (2)

Measuring time

- Wall-clock time (time-based definition)
 - All updates since 13:32:00
 - [current time – 100 seconds, current time]
 - Event time vs processing time
- Number of updates (event-based definition)
 - Last 1000 updates
 - Also called count-based, arrival-based

Two choices to handle streams

- Scale-out to handle real-time requirements
 - Add more machines
 - Programming model to ease
 - Expressing user requirements
 - Distributing the stream and computations
- The *poor man's* approach*
 - Summarize the stream
 - Approximate user requirements

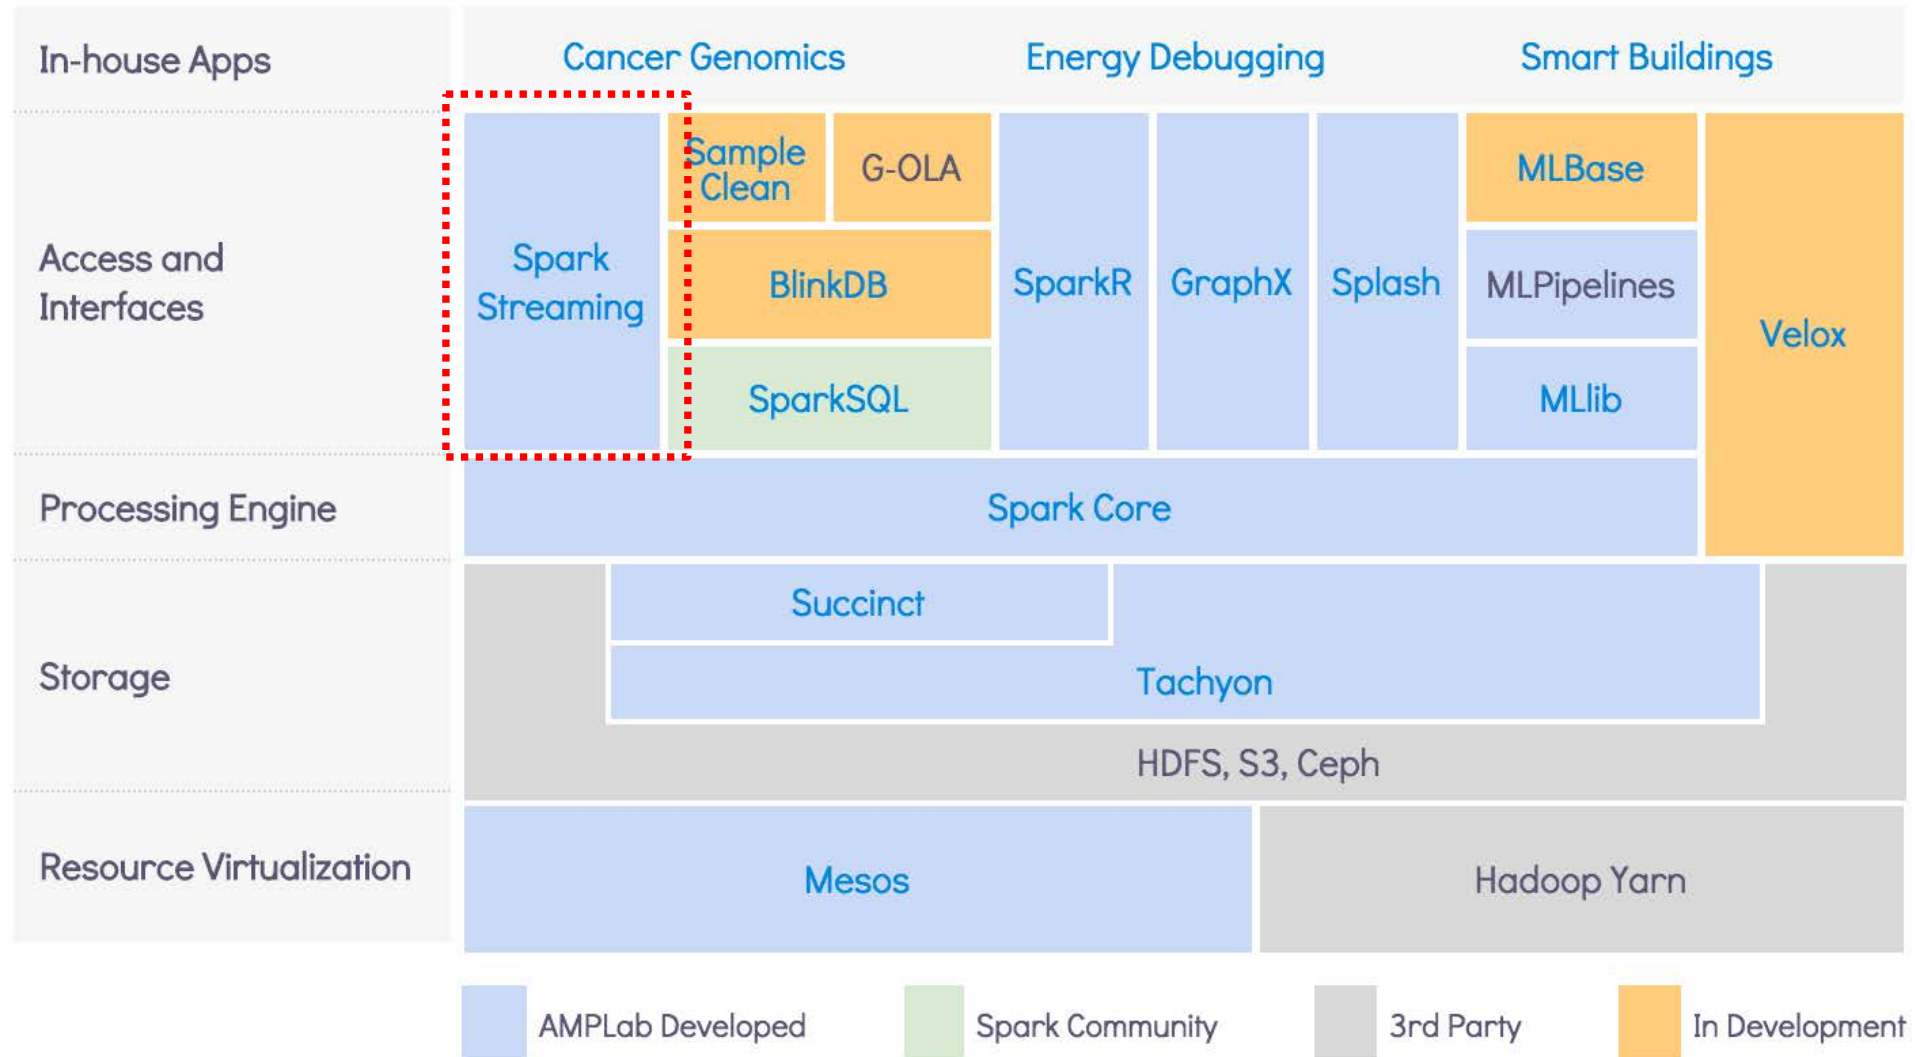
Sometimes, combination of both!

Scaling-out platforms


- Necessary when exact results are needed
 - Banks, medical sensors, Industry 4.0
- Typically comes with a hefty price tag
- Several platforms
 - Spark Streaming
 - Twitter: Storm → Heron
 - Apache Flink
 - Apache Kafka



The Spark Unified Stack



Spark Streaming

- Up to now
 - Data pre-existed in DFS
 - Analyzed in batch

One-shot queries over stored data
- Real-time analysis of big data
 - Process data *as soon as* it arrives, take action immediately
 - Continuous queries


Expectations from the platform

- Programming model
 - Utilize commodity clusters for scalable processing
- Fault tolerance
 - Must recover from **failures** and **stragglers** quickly and efficiently
 - Imagine the VISA fraud detection system is down!
- It was *almost easy* in MapReduce and Spark! Why is it so difficult for streams?

Fault tolerance for SP systems

- Replication

- Replicate processing of each input to two or more nodes
- When a node fails, the backup node takes over




Fast recovery

Very expensive!

Fault tolerance for SP systems

- Upstream backup
 - Nodes maintain checkpoints (**safely**) and backups of updates forwarded after the checkpoints
 - When a node fails, the backup node takes over
 - Replay input stream after last surviving checkpoint




Fairly cheap – minimal additional hw

Slow recovery

Fault tolerance for SP systems

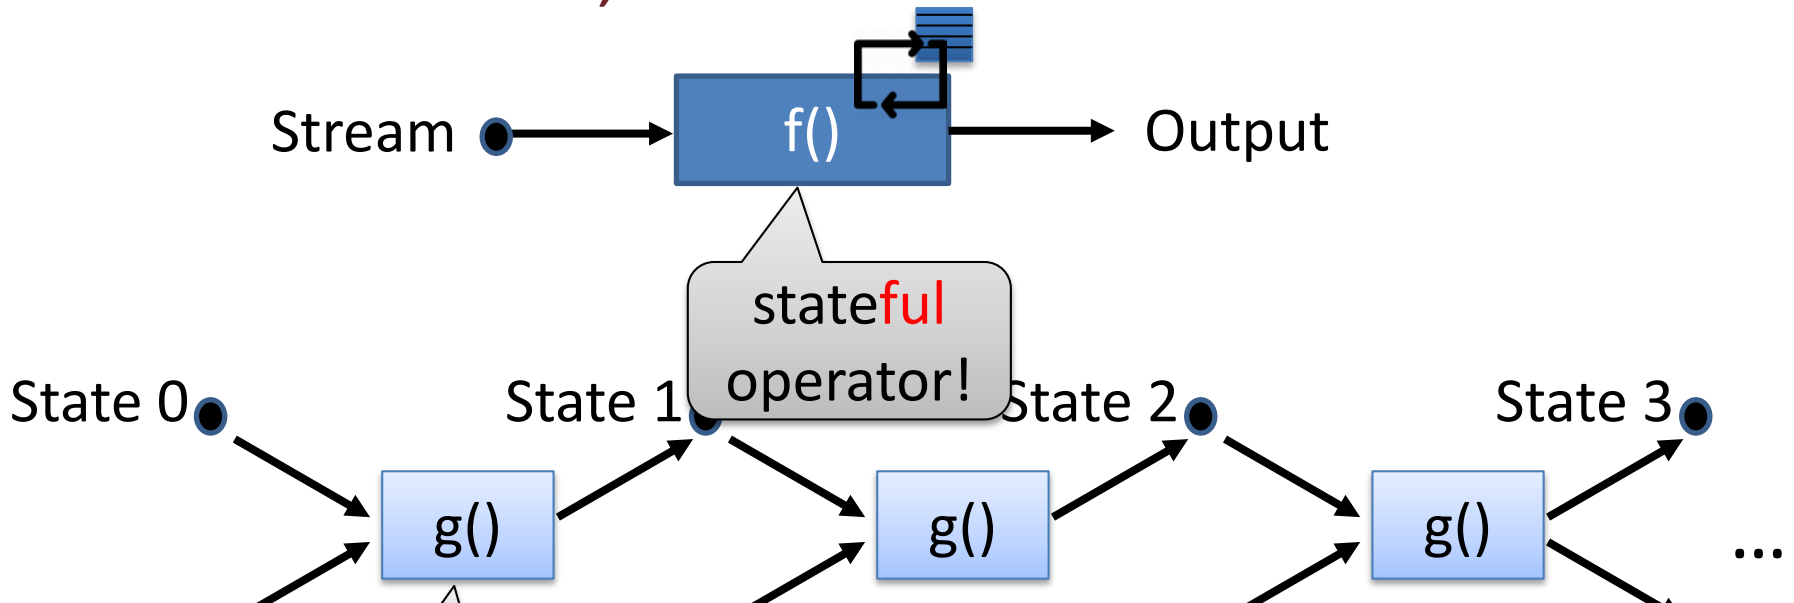
- Replication
 - Fast recovery
 - Expensive on resources – at least x2 nodes
- Upstream backup
 - Cheap – only a few additional nodes
 - Time-consuming recovery



Computation is **closely coupled** with state


A new fault tolerance technique for SP

- State: mutable \rightarrow immutable
- Partition continuous computation into **small, deterministic, stateless tasks**



We have decoupled
computation from state!

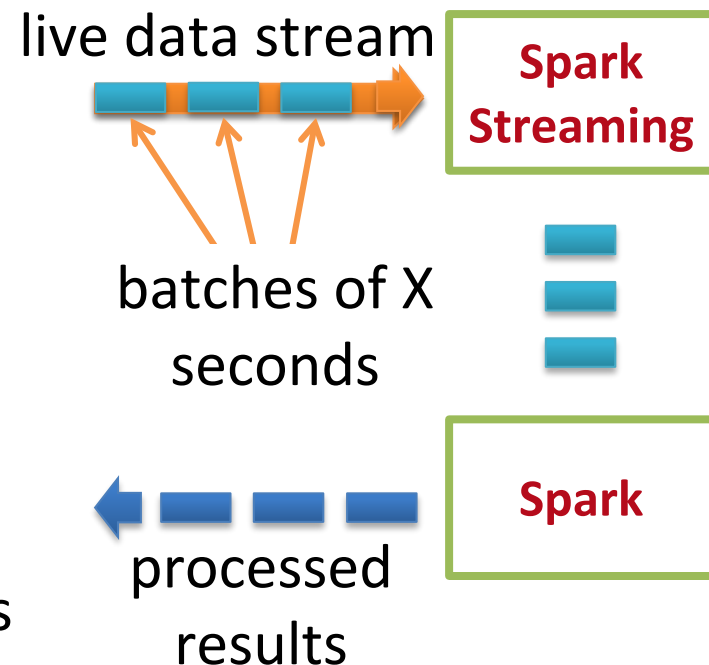
Does this ring a bell?



- **Stateless operators** → static batch model
 - MapReduce, Spark, ...
 - Already advanced in fault tolerance!

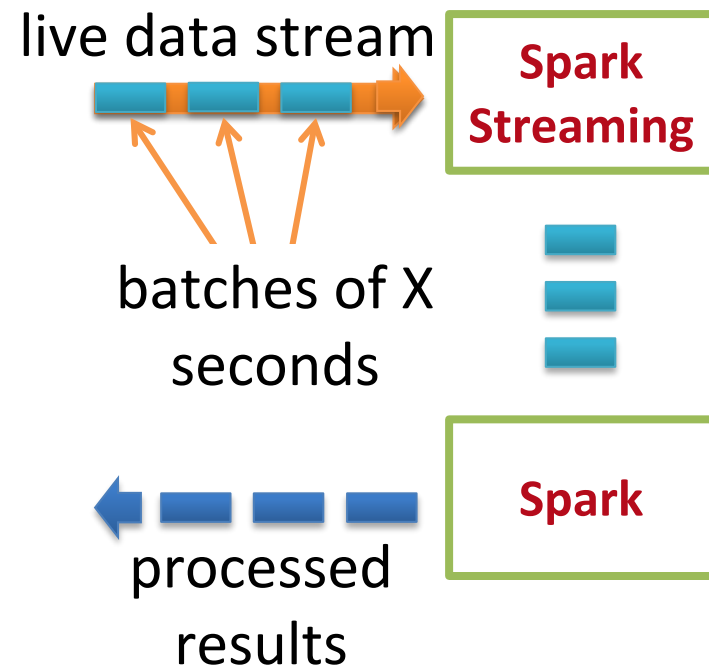
Discretized stream processing

- Consider stream in SMALL input batches
 - Micro-batches
- Streaming computation is a series of very small deterministic batch jobs
- Chop up the live stream into batches of X seconds
- Spark treats each batch of data as RDDs and processes them using RDD operations
- Finally, the processed results of the RDD operations are returned in batches



Discretized stream processing

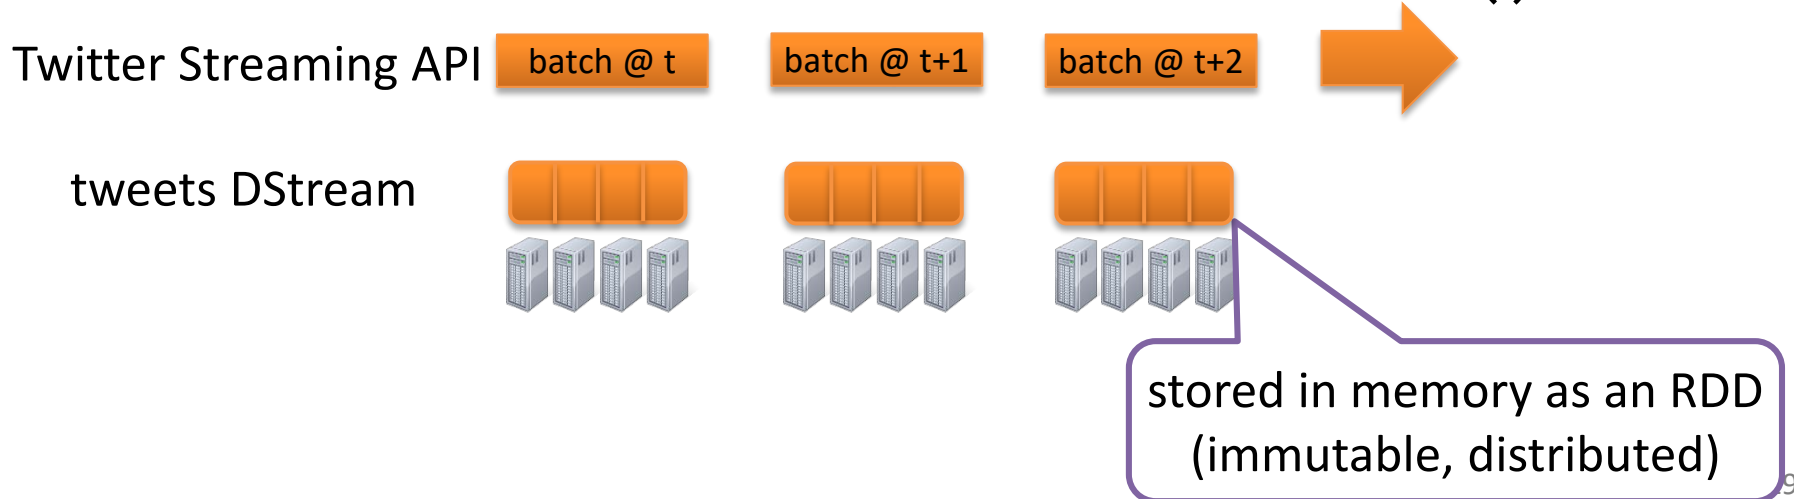
- Consider stream in SMALL input batches
 - Micro-batches
- Streaming computation is a series of very small deterministic batch jobs
- Batch sizes as low as $\frac{1}{2}$ second, latency of about 1 second
- Potential for combining batch processing and streaming processing in the same system



DStreams

- Discretized stream (DStream) is a sequence of **immutable, partitioned datasets**
 - a sequence of RDDs!
- Can be created from live data streams (e.g., twitter, network sockets, ...) or by applying bulk parallel transformations on other DStreams

```
val tweets = ssc.twitterStream()
```



Example – streaming word count

```
val ssc = new StreamingContext(..., Seconds(5))
```

streaming content

size of each batch

Twitter Streaming API

batch @ t

batch @ t+1

batch @ t+2



tweets DStream



Example – streaming word count

```
val ssc = new StreamingContext(..., Seconds(5))
val lines = ssc.socketTextStream(url, port)
```

As a DStream

read data from socket

Twitter Streaming API

batch @ t

batch @ t+1

batch @ t+2



tweets DStream



Example – streaming word count

```
val ssc = new StreamingContext(..., Seconds(5))
val lines = ssc.socketTextStream(url, port)
val words = lines.flatMap(_.split(" "))
val wordC = words.map(x => (x, 1)).reduceByKey(_ + _)
```

DStreams

transformation: modify data in one DStream to create another DStream

Twitter Streaming API

batch @ t

batch @ t+1

batch @ t+2



tweets DStream




Example – SW streaming word count

```
val ssc = new StreamingContext(..., Seconds(5))
val lines = ssc.socketTextStream(url, port)
val words = lines.window(Seconds(10))
               .flatMap(_.split(" "))
               .map(x => (x, 1)).reduceByKey(_ + _)

words.print()
ssc.start()
ssc.awaitTermination()
```

operate on sliding
window of 10
seconds → last 2
batches



Spark Streaming *magically* maintains
sliding window operations

Example – SW streaming word count

- Defining the sliding window

Using window

```
lines.window(Seconds(10))  
      .flatMap(_.split(" "))  
      .map(x => (x, 1)).reduceByKey(_ + _)
```

Using reduceByKeyAndWindow

```
lines.flatMap(_.split(" ")).map(x => (x,1))  
      .reduceByKeyAndWindow( (a,b) => (a + b), Seconds(10) )
```

... and teaching Spark how to expire batches

```
lines.flatMap(_.split(" ")).map(x => (x,1))
```

What is the benefit of the last approach?

Arbitrary Combinations of Batch and Streaming Computations

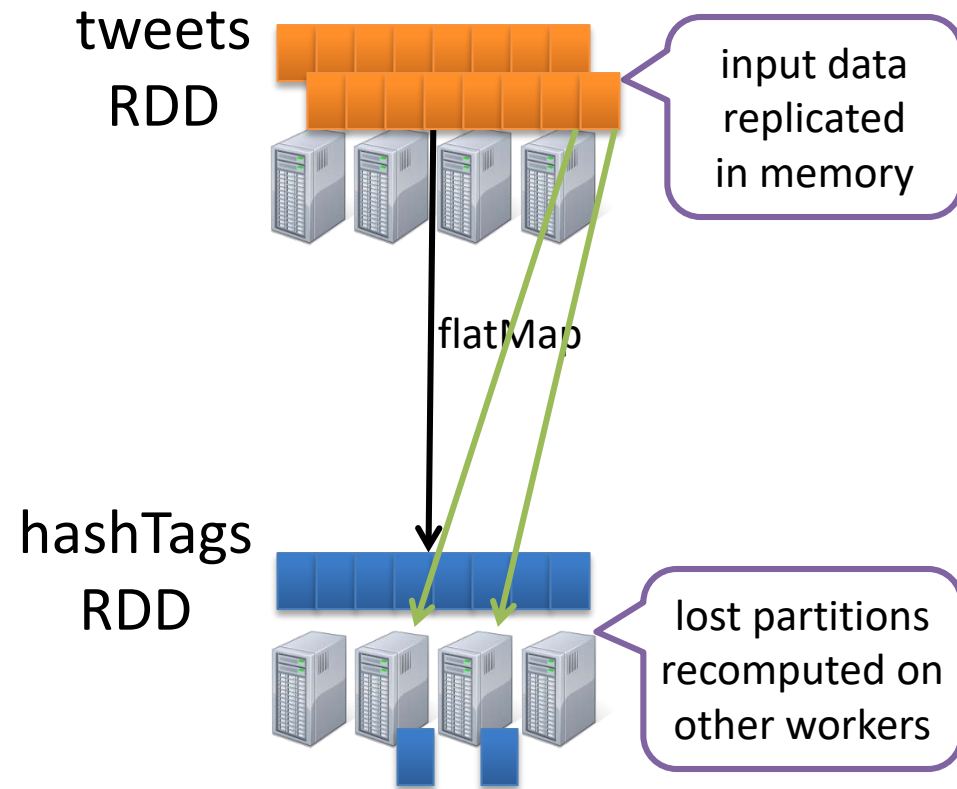
Inter-mix RDD and DStream operations!

Example: Join incoming tweets with a spam filter for data cleaning

```
tweets.transform(tweetsRDD => {  
    tweetsRDD.join(spamHDFSFile).filter(...)  
})
```


Fault-tolerance

- RDD Lineage: RDDs remember the operations that created them
- Batches of input data are replicated in memory for fault-tolerance
- Data lost due to worker failure, can be recomputed from replicated input data
- Fault-tolerance and exactly-once transformations



Fast Fault Recovery

Recovers from faults/stragglers within **1 sec**



Unifying Batch and Stream Processing Models

Spark program on Twitter log file using RDDs

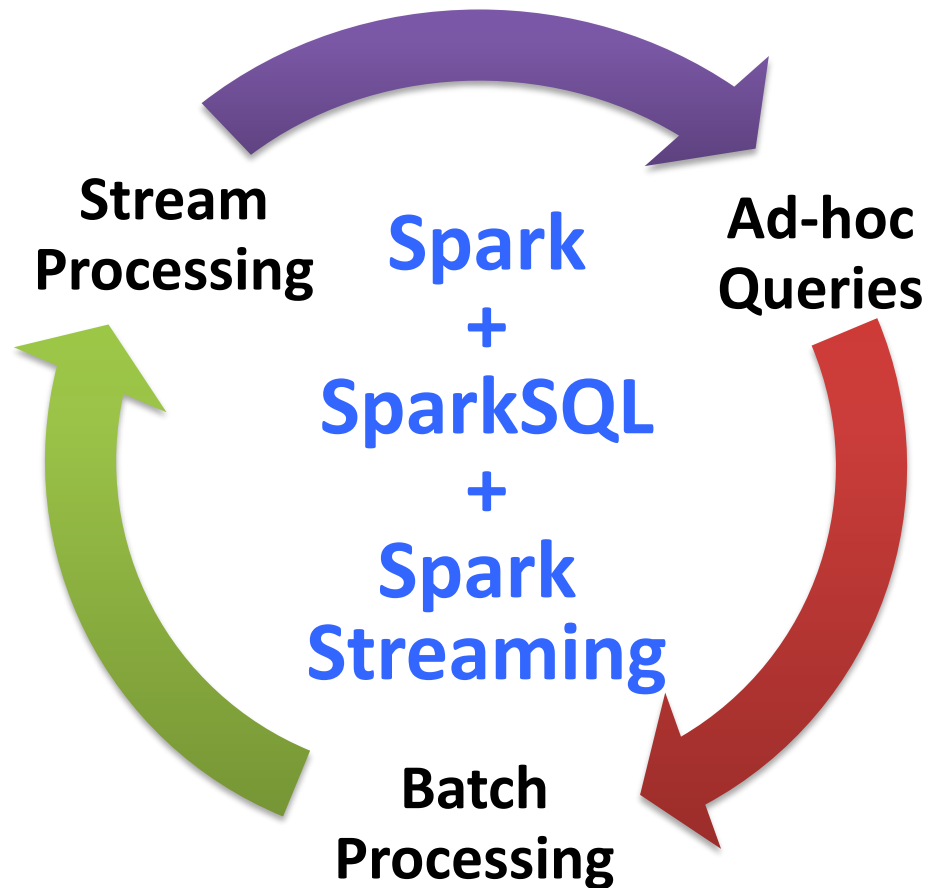
```
val tweets = sc.hadoopFile("hdfs://...")  
val hashTags = tweets.flatMap(status => getTags(status))  
hashTags.saveAsHadoopFile('hdfs://...')
```

ONLY NEED TO CHANGE THE INPUT!

Spark Streaming program on Twitter stream using DStreams

```
val tweets = ssc.twitterStream()  
val hashTags = tweets.flatMap(status => getTags(status))  
hashTags.saveAsHadoopFiles("hdfs://...")
```

Vision - *one stack to rule them all*



OTHERSIDE

RED HOT CHILI PEPPERS



Our progress up to now!

100 TB on 1000 machines

½ - 1 Hour



Hard Disks

1 - 5 Minutes



Memory

1-5 seconds



Approximate Query Processing

Exact

Approximate

Offline



Real-time



Synopses &
approx. alg.

An example...

- Server log files of several TB

```
64.242.88.10 - - [07/Mar/2004:16:05:49 -0800] "GET /twiki/bin/edit/Main/Double_...
142.22.55.12 - - [07/Mar/2004:16:06:51 -0800] "GET /twiki/bin/rdiff/TWiki/NewUs...
142.242.63.63 - - [07/Mar/2004:16:10:02 -0800] "GET /mailman/listinfo/hsdivision...
142.22.55.12 - - [07/Mar/2004:16:11:58 -0800] "GET /twiki/bin/view/TWiki/WikiSy...
111.11.32.65 - - [07/Mar/2004:16:20:55 -0800] "GET /twiki/bin/view/Main/DCCAndP...
112.13.45.99 - - [07/Mar/2004:16:23:12 -0800] "GET /twiki/bin/oops/TWiki/Append...
62.11.1.123 - - [07/Mar/2004:16:24:16 -0800] "GET /twiki/bin/view/Main/PeterTh...
```

...

- How to query these files
 - using good old **SQL**
 - and get **near-interactive responses?**
 - **approximate results** acceptable!!!

BlinkDB: Blink and it's done

- Target: Support **interactive** SQL-like aggregate queries over massive sets of data by **approximation**
- Interactive: ~ 1-2 seconds
- Strategy: Approximate answers via sampling
- Supports
 - Aggregates & group-by [Paper 1, in exam]
 - Filters [Paper 1, in exam]
 - Joins [Paper 2, **not in exam material**]
 - User-defined functions (UDFs)! [Paper 2, **not in exam material**]

Paper 1 (in exam): https://sameeragarwal.github.io/blinkdb_eurosys13.pdf

Paper 2 (optional, will not be examined in midterm or final):

<https://sameeragarwal.github.io/mod282-agarwal.pdf>

BlinkDB (2)

Supports

Aggregates

```
•blinkdb> SELECT AVG( jobtime )  
FROM very_big_log
```

AVG, COUNT, SUM,
STDEV, PERCENTILE
etc.

BlinkDB (2)

Supports

Aggregates

Filters

```
•blinkdb> SELECT AVG( jobtime )  
FROM very_big_log  
WHERE src = 'hadoop'
```



FILTERS, GROUP BY clauses

BlinkDB (2)

Supports

Aggregates

Filters

Joins

```
•blinkdb> SELECT AVG( jobtime )  
FROM very_big_log  
WHERE src = 'hadoop'
```

```
LEFT OUTER JOIN logs2  
ON very_big_log.id = logs.id
```

JOINS, Nested Queries etc.

BlinkDB (2)

Supports

Aggregates
Joins

Filters
UDFs

• `blinkdb>`

```
SELECT my_function(jobtime)
```

```
FROM very_big_log
```

```
WHERE src = 'hadoop'
```

```
LEFT OUTER JOIN logs2
```

```
ON very_big_log.id = logs.id
```

ML Primitives,
User Defined Functions

BlinkDB (2)

Supports

Aggregates

Filters

Joins

UDFs

Accuracy reqs

```
•blinkdb> SELECT AVG(jobtime)
FROM very_big_log
WHERE src = 'hadoop'
LEFT OUTER JOIN logs2
ON very_big_log.id = logs.id
```

ERROR WITHIN 10% AT CONFIDENCE 95%

Desired accuracy

BlinkDB (2)

Supports

Aggregates

Filters

Joins

UDFs

Accuracy reqs

Performance reqs

```
•blinkdb> SELECT AVG(jobtime)
FROM very_big_log
WHERE src = 'hadoop'
LEFT OUTER JOIN logs2
ON very_big_log.id = logs.id
WITHIN 5 SECONDS
```

Desired performance

Workflow of BlinkDB

ID	City	Buff Ratio
1	NYC	0.78
2	NYC	0.13
3	Berkeley	0.25
4	NYC	0.19
5	NYC	0.11
6	Berkeley	0.09
7	NYC	0.18
8	NYC	0.15
9	Berkeley	0.13
10	Berkeley	0.49
11	NYC	0.19
12	Berkeley	0.10

Query log analysis:

Q1: SELECT count(*) FROM log WHERE city=NYC

Q2: SELECT AVG(bradio) FROM log WHERE city=NYC

Q3: SELECT AVG(bradio) FROM log GROUP BY city

...

Decide on
the samples



...

ID	City	Buff Ratio
1	City	Buff Ratio
3	NYC	0.25
4	NYC	0.19
9	NYC	0.15
10	Berkeley	0.78
10	Berkeley	0.13
12	Berkeley	0.19
	Berkeley	0.09

Query Execution on Samples

ID	City	Buff Ratio
1	NYC	0.78
2	NYC	0.13
3	Berkeley	0.25
4	NYC	0.19
5	NYC	0.11
6	Berkeley	0.09
7	NYC	0.18
8	NYC	0.15
9	Berkeley	0.13
10	Berkeley	0.49
11	NYC	0.19
12	Berkeley	0.10

What is the average buffering ratio in the table?

0.2325

Query Execution on Samples

ID	City	Buff Ratio
1	NYC	0.78
2	NYC	0.13
3	Berkeley	0.25
4	NYC	0.19
5	NYC	0.11
6	Berkeley	0.09
7	NYC	0.18
8	NYC	0.15
9	Berkeley	0.13
10	Berkeley	0.49
11	NYC	0.19
12	Berkeley	0.10

→
Uniform
Sample

What is the average buffering ratio in the table?

ID	City	Buff Ratio	Sampling Rate
2	NYC	0.13	1/4
6	Berkeley	0.25	1/4
8	NYC	0.19	1/4

~~0.2325~~

0.19 +/- 0.05

Query Execution on Samples

ID	City	Buff Ratio
1	NYC	0.78
2	NYC	0.13
3	Berkeley	0.25
4	NYC	0.19
5	NYC	0.11
6	Berkeley	0.09
7	NYC	0.18
8	NYC	0.15
9	Berkeley	0.13
10	Berkeley	0.49
11	NYC	0.19
12	Berkeley	0.10

→
Uniform
Sample

What is the average buffering ratio in the table?

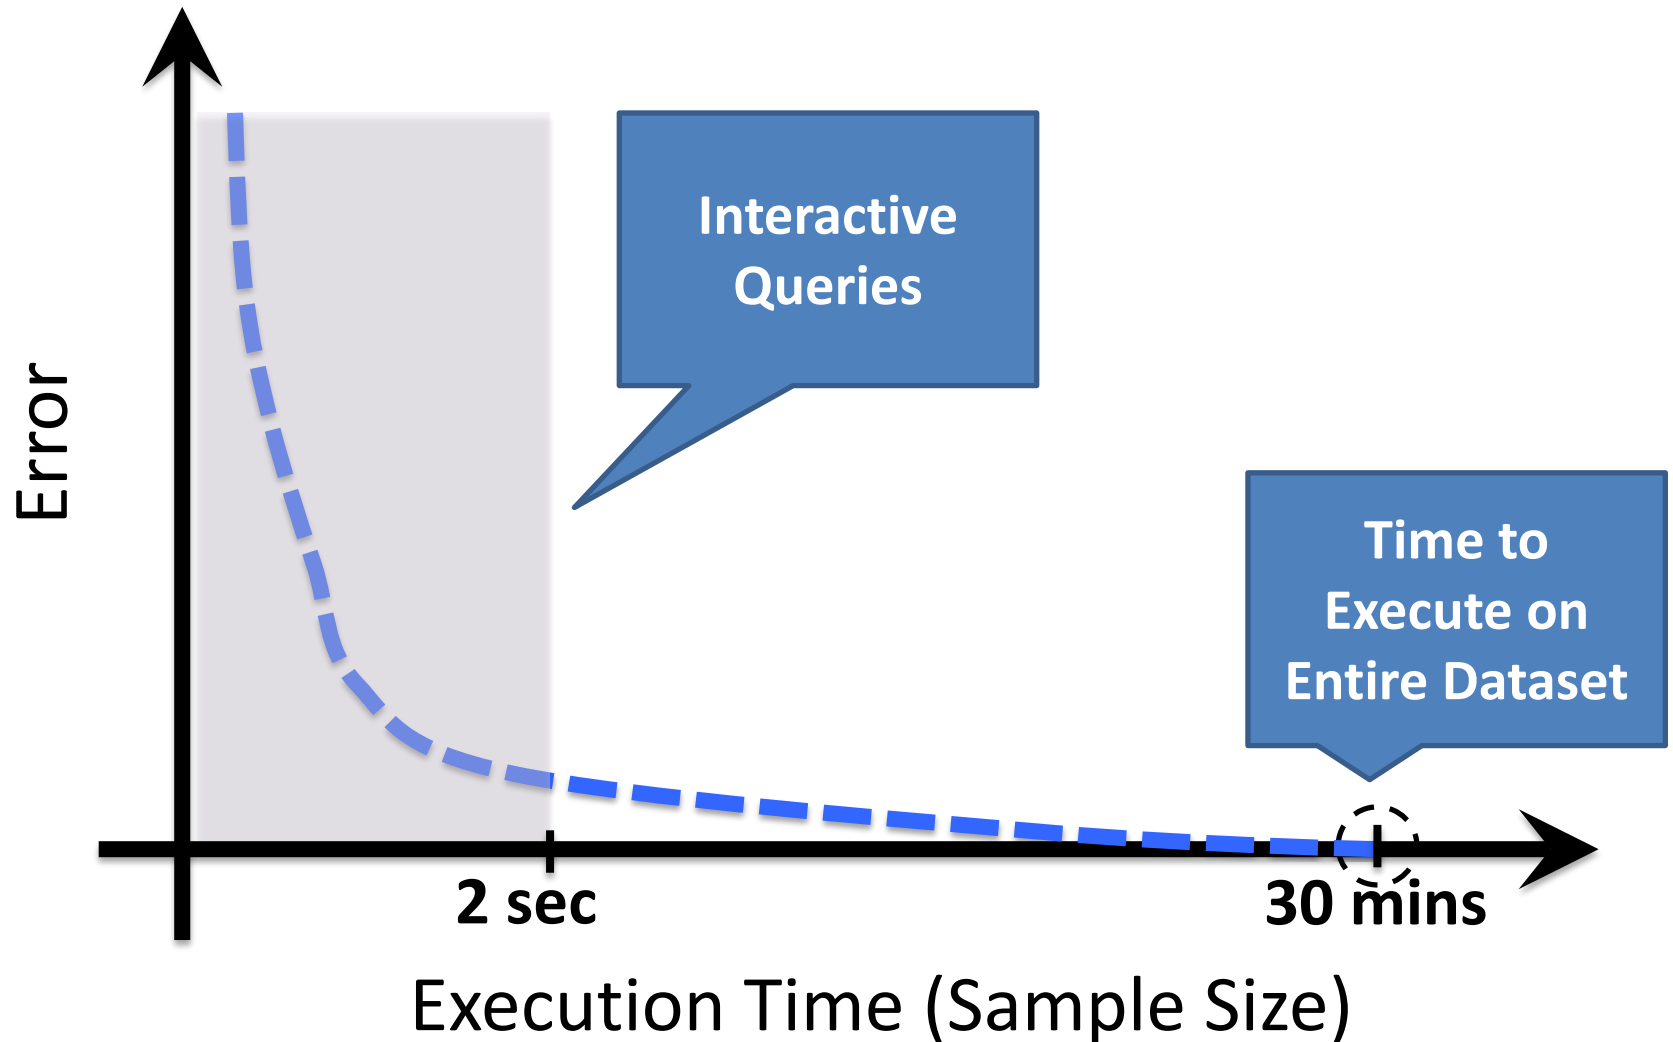
ID	City	Buff Ratio	Sampling Rate
2	NYC	0.13	1/2
3	Berkeley	0.25	1/2
5	NYC	0.19	1/2
6	Berkeley	0.09	1/2
8	NYC	0.18	1/2
12	Berkeley	0.49	1/2

~~0.2325~~


~~0.19~~ +/- ~~0.05~~

0.22 +/- 0.02

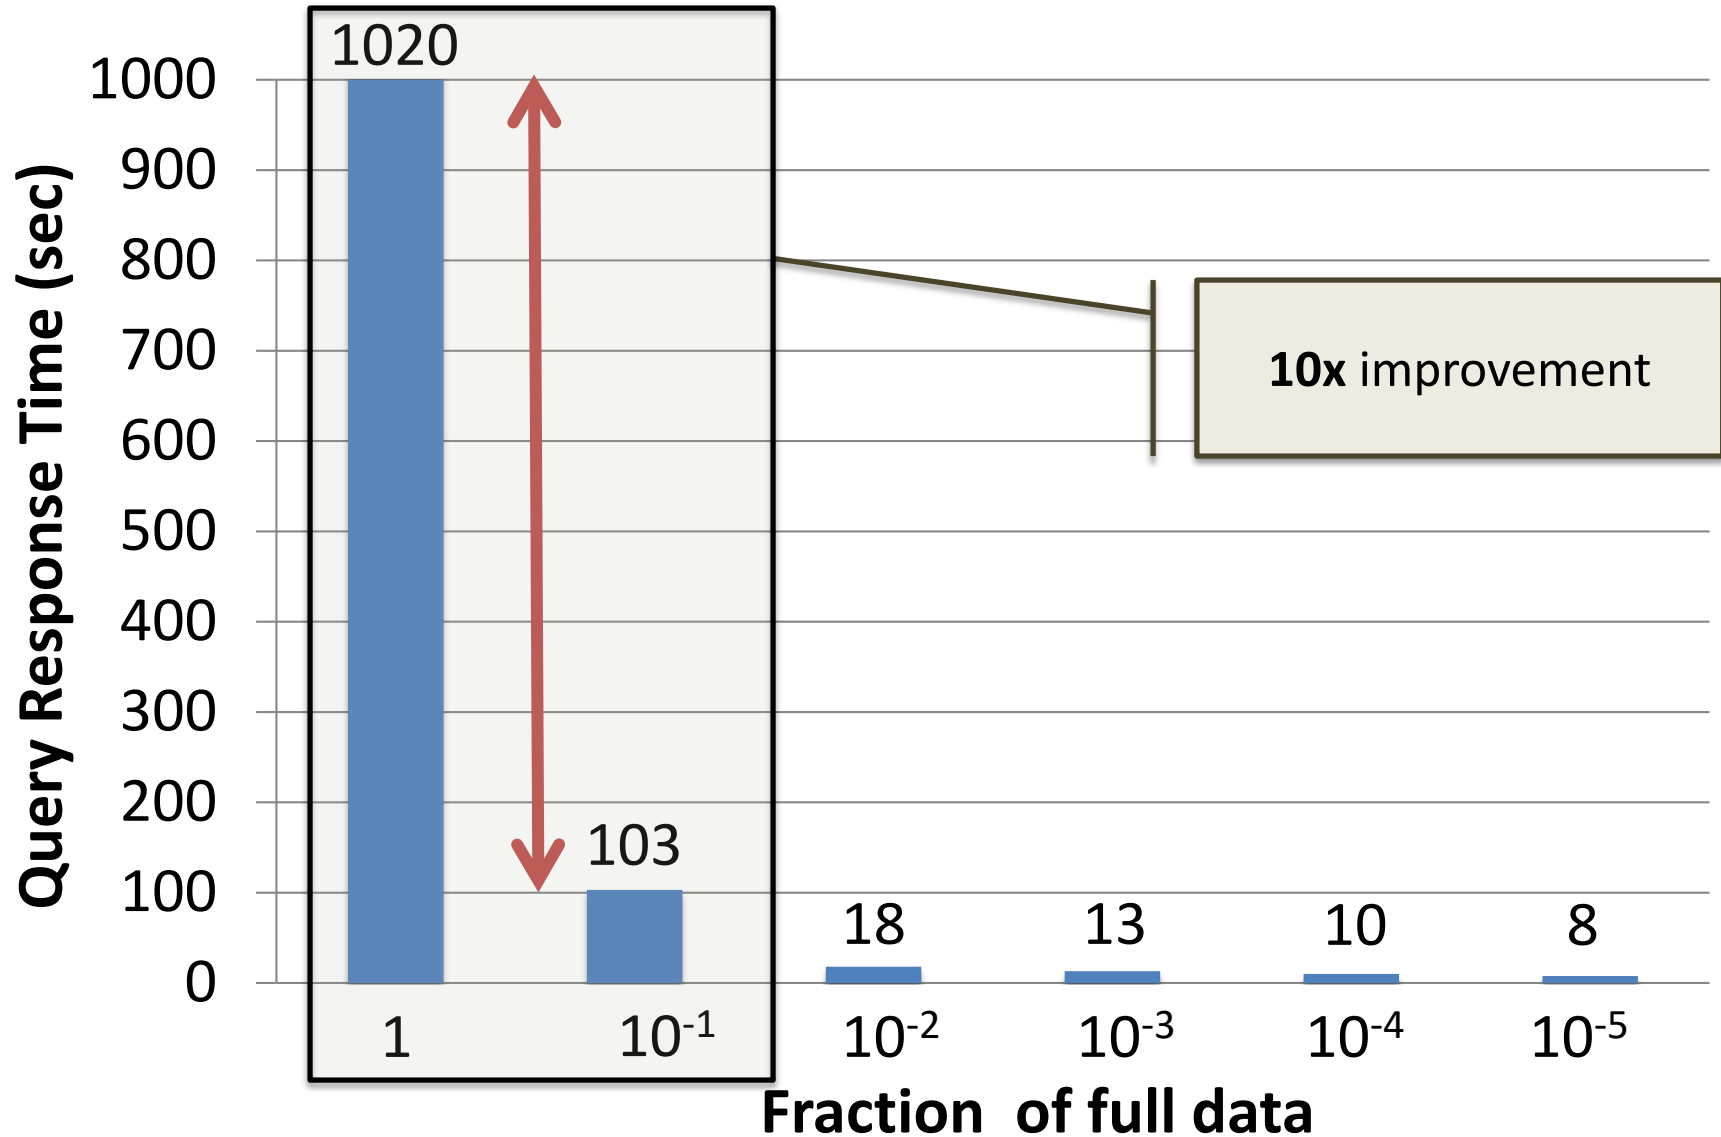
Speed/Accuracy Trade-off



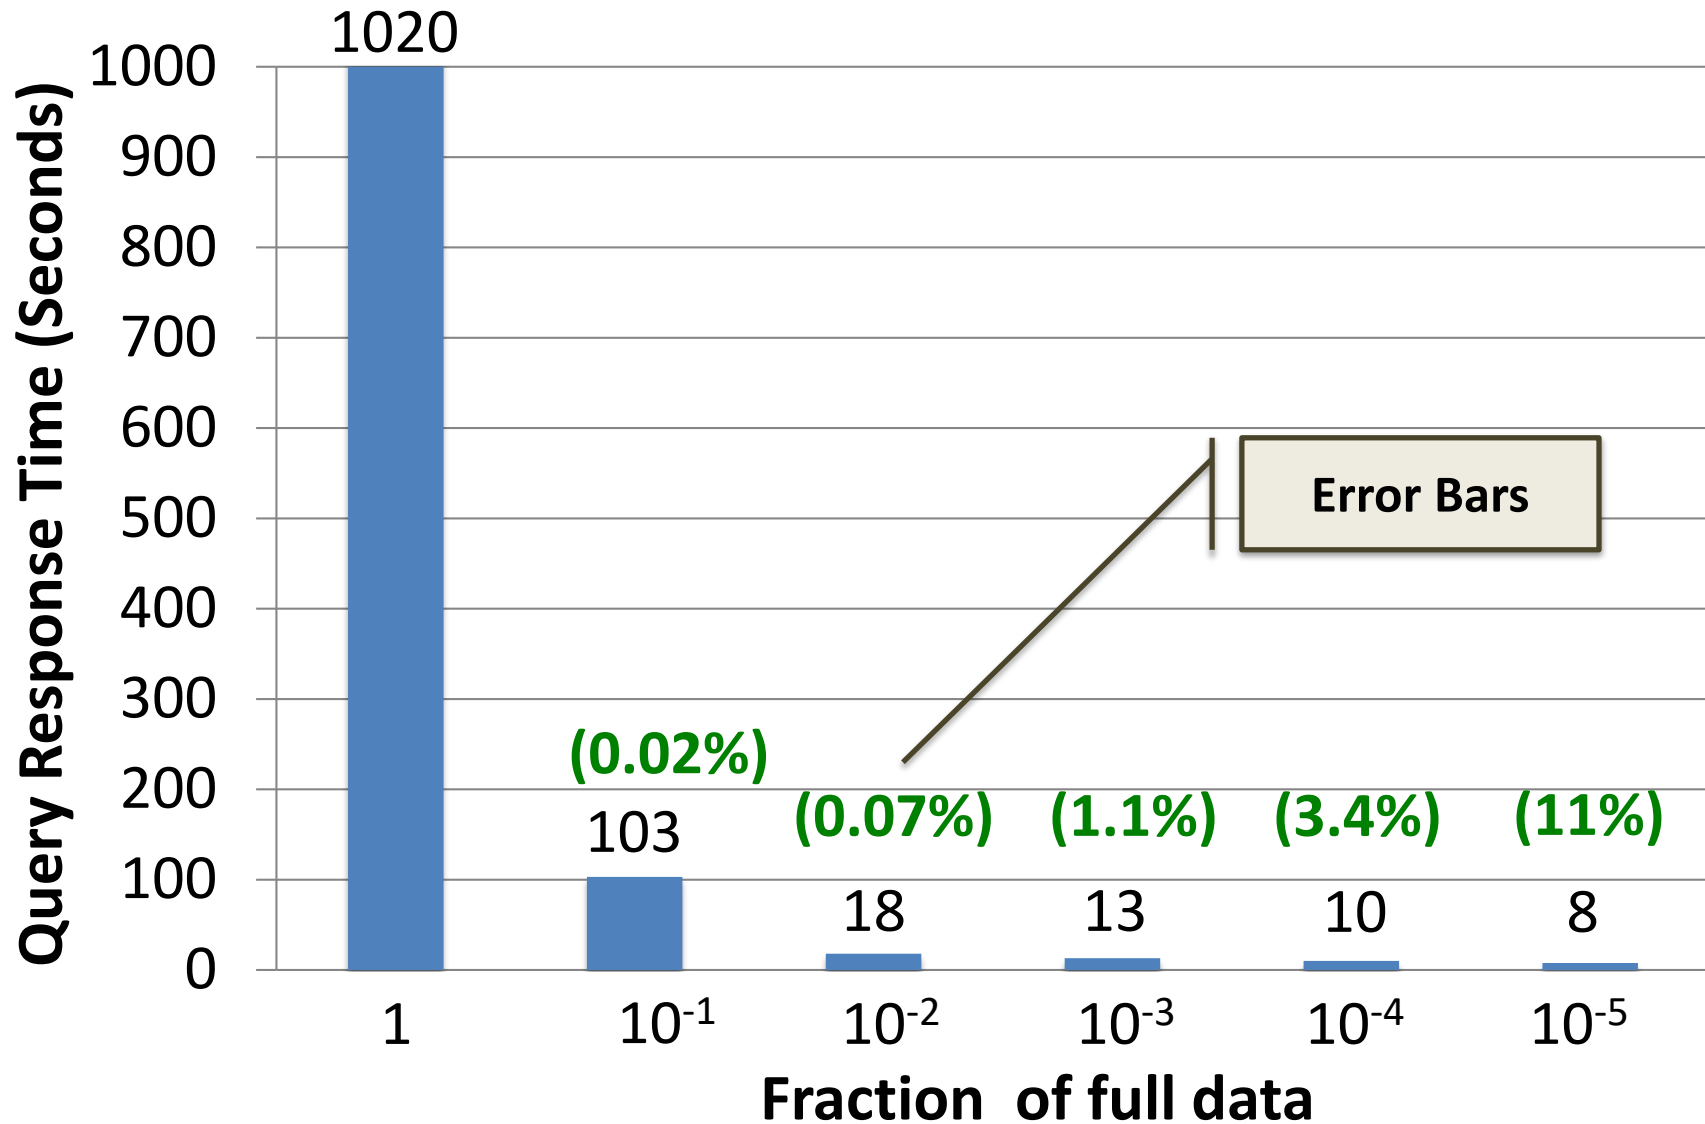
Speed/Accuracy Trade-off



Sampling Vs. No Sampling



Sampling Vs. No Sampling



What is BlinkDB?

- A framework that ...
 - creates and maintains a variety of **uniform** and **stratified** samples from underlying data
 - returns fast, approximate answers with error bars by executing queries on samples of data
 - verifies the correctness of the error bars that it returns at runtime

What is BlinkDB?

- A framework that ...
 - creates and maintains a variety of uniform and stratified samples from underlying data
 - returns fast, approximate answers with error bars by executing queries on samples of data
 - verifies the correctness of the error bars that it returns at runtime

Learning to sample!

- Which types of sample to create?
 - On which columns
 - How many samples
- Typical assumption
 - Interests don't change → let the past queries guide you!

```
SELECT AVG(salary)
FROM tbl1
```

Uniform sample

```
SELECT AVG(salary)
FROM tbl1
WHERE city="London"
```

Uniform sample
for city="London"



```
SELECT
AVG(salary)
FROM tbl1
GROUP BY city,
        profession
```

???

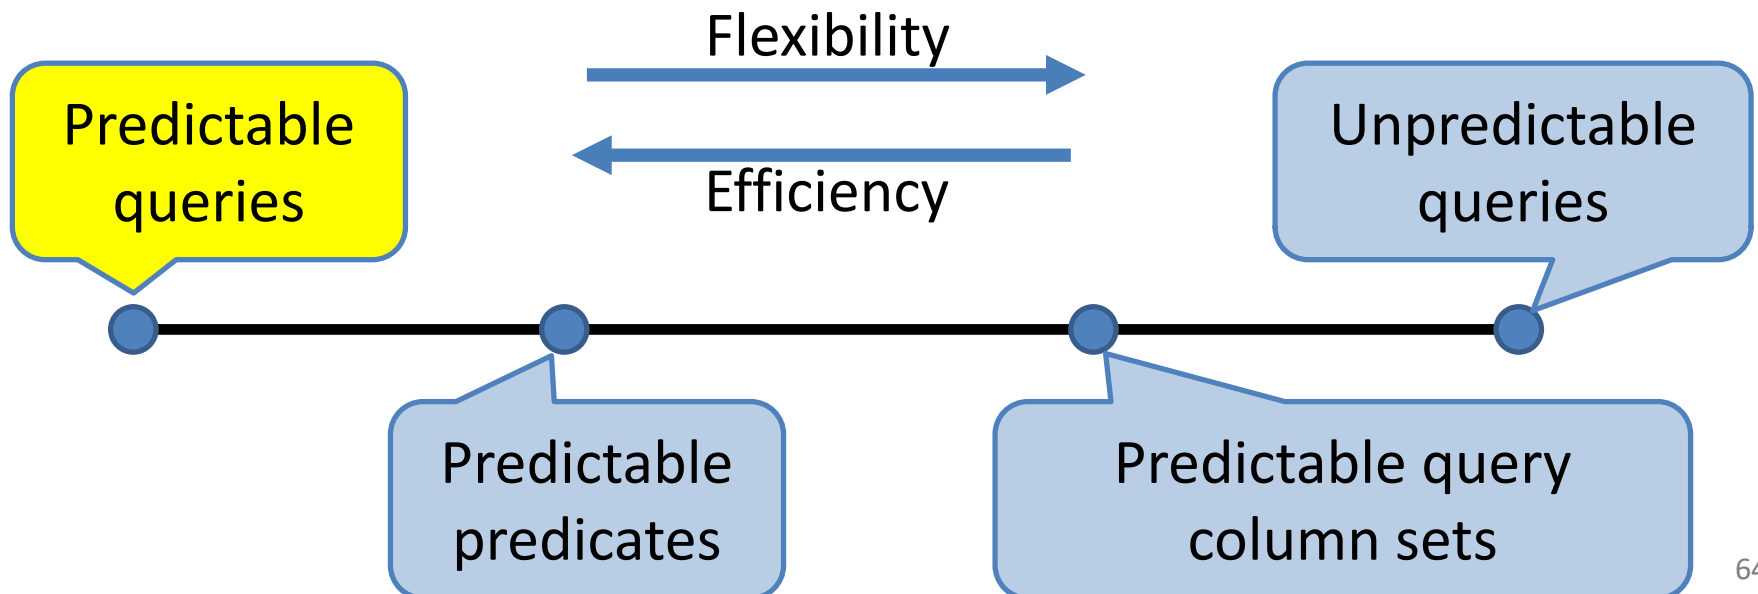
???

Learning to sample (2)

- Predictable queries

```
SELECT AVG(salary)FROM tbl1
WHERE city="London"
```

- Can only answer this exact query!

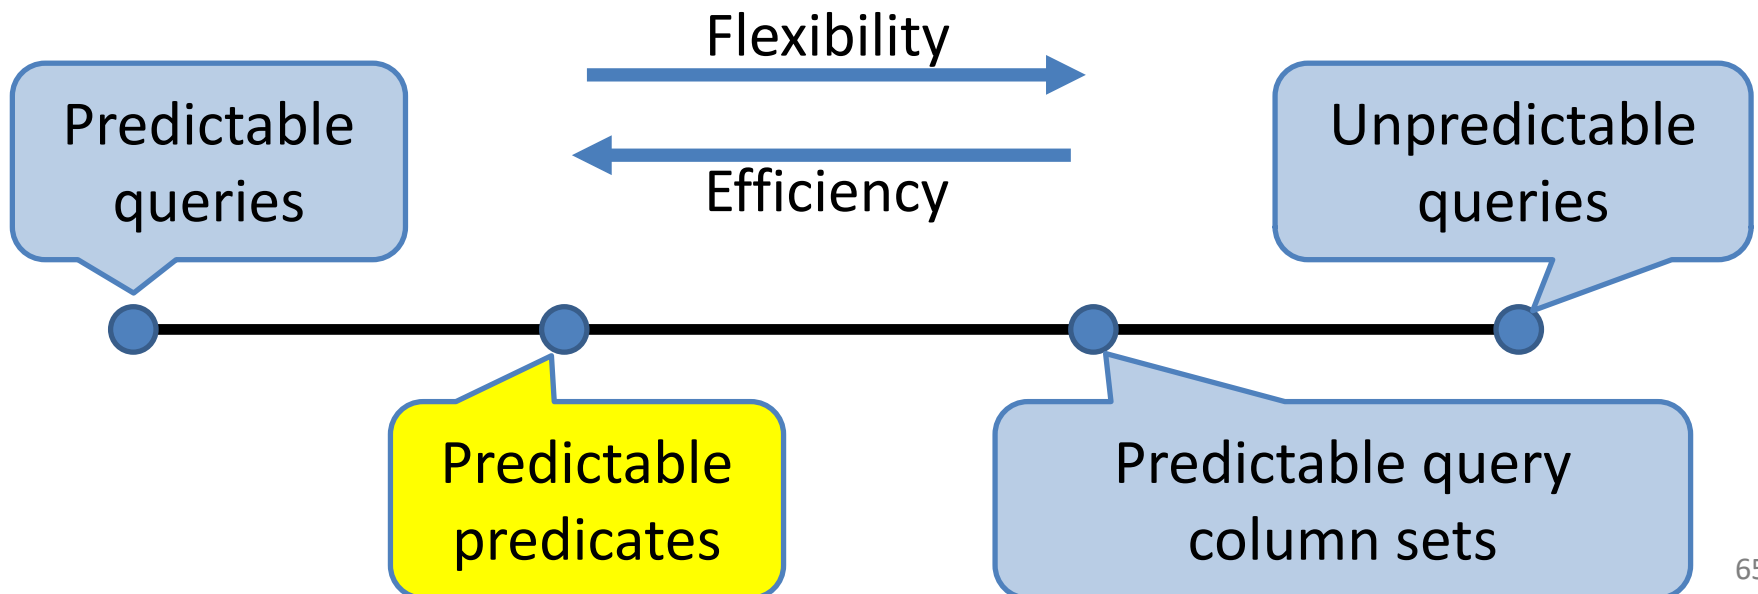


Learning to sample (3)

- Predictable predicates

```
SELECT AVG($X) FROM tbl1
WHERE city="London"
```

- Can only answer this query, for different \$X

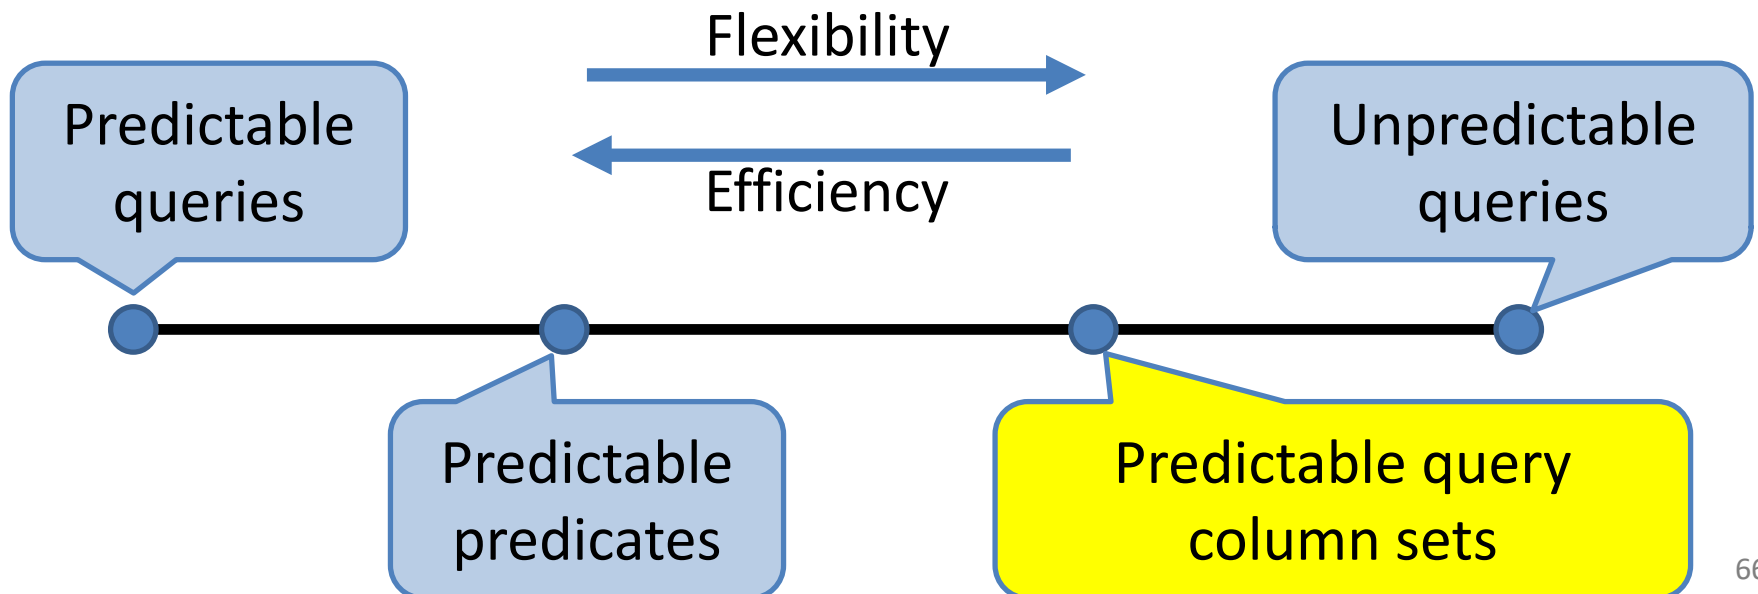


Learning to sample (4)

- Key notion: QCS - Query Column Sets
 - All **columns** contained in the query that affect sampling
 - **GROUP BY, HAVING, WHERE**

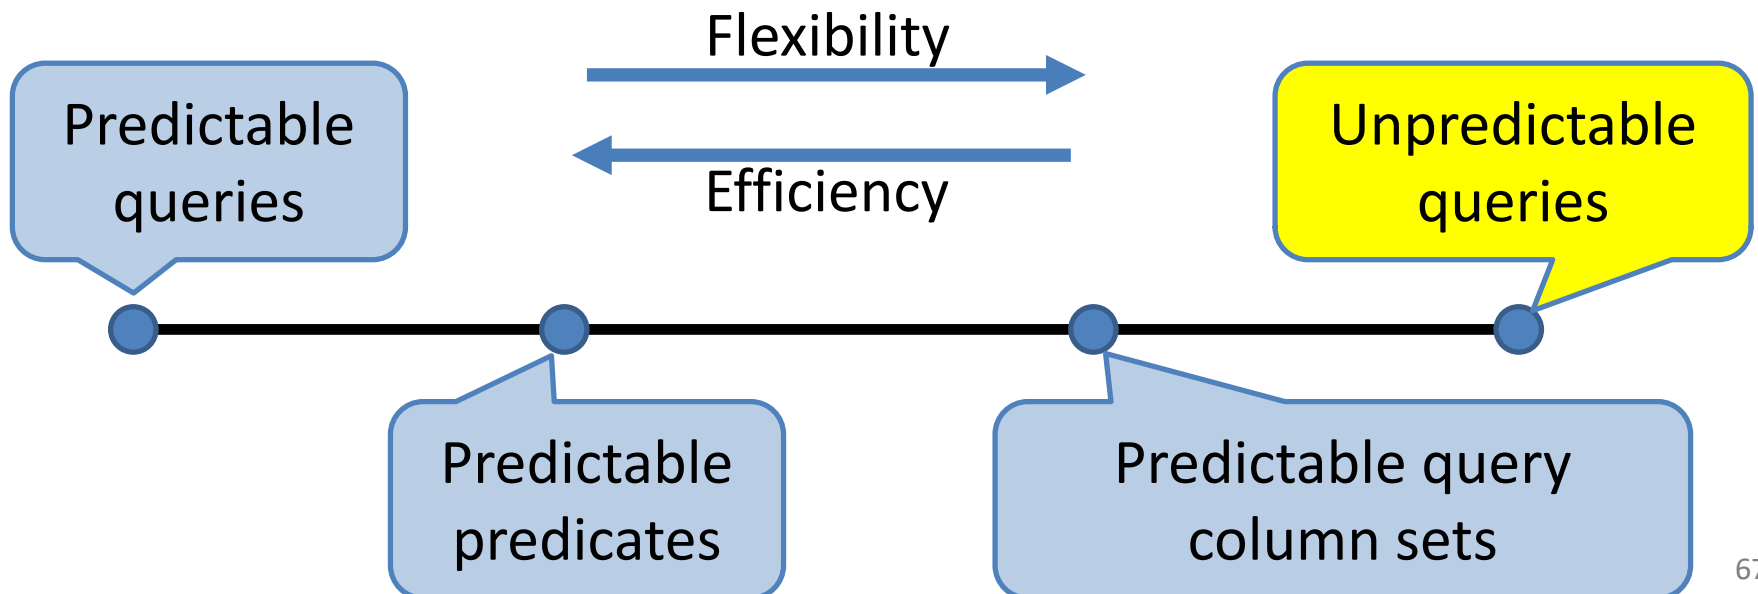
```
SELECT AVG($X) FROM tbl1
WHERE $Y1=$val1 AND $Y2=$val2 ... GROUP BY $Z
```

- Construct samples for each QCS



Learning to sample (5)

- Unpredictable queries
- Fully flexible, but cannot sample efficiently!
 - Best effort approach



Query column sets

- BlinkDB uses QCS to sample
 - Over 90% of queries covered by 10% - 20% of QCS →
 - Keeping samples for the most frequently used 10%-20% of QCS can help in 90% of the queries!
 - Use query logs to find the most frequent QCS
- **Maintain a different sample for each frequent QCS!**
 - For each query q with QCS_q
 - if there exists already a suitable sample S with QCS_S , use it
 - Suitable: QCS_q is a subset of QCS_S
 - If there is no suitable sample, trial-and-error!

Query column sets (2)

- Maintain a different sample for each frequent QCS!

Pair the queries with the samples

- Samples on the following QCS

- <city>
- <city, profession>
- <city, age>
- <city, profession, age>
- <city, profession, education>
- <education>

Queries

```
SELECT avg(salary)
FROM tbl1
GROUP BY
city, profession
```

Query column sets (2)

- Maintain a different sample for each frequent QCS!

Pair the queries with the samples

- Samples on the following QCS

- <city>
- <city, profession>
- <city, age>
- <city, profession, age>
- <city, profession, education>
- <education>

Queries

```
SELECT avg(salary)
FROM tbl1
WHERE age=20
GROUP BY
city
```

Query column sets (2)

- Maintain a different sample for each frequent QCS!

Pair the queries with the samples

- Samples on the following QCS

- <city>
- <city, profession>
- <city, age>
- <city, profession, age>
- <city, profession, education>
- <education>

Queries

```
SELECT avg(salary)
FROM tbl1
WHERE city=London
GROUP BY
education
```

Query column sets (2)

- Maintain a different sample for each frequent QCS!

Pair the queries with the samples

- Samples on the following QCS

- <city>
- <city, profession>
- <city, age>
- <city, profession, age>
- <city, profession, education>
- <education>

Queries

```
SELECT avg(salary)
FROM tbl1
GROUP BY
city, education
```

Query column sets (2)

- Maintain a different sample for each frequent QCS!

Pair the queries with the samples

- Samples on the following QCS


- <city>
- <city, profession>
- <city, age>
- <city, profession, age>
- <city, profession, education>
- <education>

Queries

```
SELECT avg(salary)
FROM tbl1
GROUP BY
name
```

Uniform samples

ID	City	Data
1	NYC	0.78
2	NYC	0.13
3	Berkeley	0.25
4	NYC	0.19
5	NYC	0.11
6	Berkeley	0.09
7	NYC	0.18
8	NYC	0.15
9	Berkeley	0.13
10	Berkeley	0.49
11	NYC	0.19
12	Berkeley	0.10



Simplification


SELECT AVG(data)
FROM tbl
GROUP BY city



SELECT AVG(data)
FROM tbl1, tbl2
WHERE tbl1.x=tbl2.y
GROUP BY city

Uniform samples

ID	City	Data
1	NYC	0.78
2	NYC	0.13
3	Berkeley	0.25
4	NYC	0.19
5	NYC	0.11
6	Berkeley	0.09
7	NYC	0.18
8	NYC	0.15
9	Berkeley	0.13
10	Berkeley	0.49
11	NYC	0.19
12	Berkeley	0.10




```

SELECT AVG(data)
FROM tbl
GROUP BY city
  
```

Uniform samples

ID	City	Data
1	NYC	0.78
2	NYC	0.13
3	Berkeley	0.25
4	NYC	0.19
5	NYC	0.11
6	Berkeley	0.09
7	NYC	0.18
8	NYC	0.15
9	Berkeley	0.13
10	Berkeley	0.49
11	NYC	0.19
12	Berkeley	0.10




1. FILTER **rand()** < 1/3
2. Adds per-row sampling rates

```
SELECT AVG(data)
FROM tbl
GROUP BY city
```

Uniform samples

ID	City	Data
1	NYC	0.78
2	NYC	0.13
3	Berkeley	0.25
4	NYC	0.19
5	NYC	0.11
6	Berkeley	0.09
7	NYC	0.18
8	NYC	0.15
9	Berkeley	0.13
10	Berkeley	0.49
11	NYC	0.19
12	Berkeley	0.10



ID	City	Data	Rate
2	NYC	0.13	1/3
8	NYC	0.25	1/3
6	Berkeley	0.09	1/3
11	NYC	0.19	1/3


Does not change
query semantics

```

SELECT AVG(data)
FROM tbl
GROUP BY city
  
```

Uniform samples

ID	City	Data
1	Lausanne	0.91
2	NYC	0.13
3	Berkeley	0.25
4	NYC	0.19
5	NYC	0.11
6	Berkeley	0.09
7	NYC	0.18
8	NYC	0.15
9	Berkeley	0.13
10	Berkeley	0.49
11	NYC	0.19
12	Berkeley	0.10



ID	City	Data	Rate
2	NYC	0.13	1/3
8	NYC	0.25	1/3
6	Berkeley	0.09	1/3
11	NYC	0.19	1/3


Uniform samples
fail to include **rare**
values

```

SELECT AVG(data)
FROM tbl
GROUP BY city
  
```

Stratified samples


ID	City	Data
1	NYC	0.78
2	NYC	0.13
3	Berkeley	0.25
4	NYC	0.19
5	NYC	0.11
6	Berkeley	0.09
7	NYC	0.18
8	NYC	0.15
9	Berkeley	0.13
10	Berkeley	0.49
11	NYC	0.19
12	Berkeley	0.10



```
SELECT AVG(data)
FROM tbl
GROUP BY city
```

Stratified samples


ID	City	Data
1	NYC	0.78
2	NYC	0.13
3	Berkeley	0.25
4	NYC	0.19
5	NYC	0.11
6	Berkeley	0.09
7	NYC	0.18
8	NYC	0.15
9	Berkeley	0.13
10	Berkeley	0.49
11	NYC	0.19
12	Berkeley	0.10



```
SELECT AVG(data)
FROM tbl
GROUP BY city
```

Stratified samples

```
SELECT AVG(data)
FROM tbl
GROUP BY city
```




ID	City	Data
1	NYC	0.78
2	NYC	0.13
3	Berkeley	0.25
4	NYC	0.19
5	NYC	0.11
6	Berkeley	0.09
7	NYC	0.18
8	NYC	0.15
9	Berkeley	0.13
10	Berkeley	0.49
11	NYC	0.19
12	Berkeley	0.10

Stratified samples

```
SELECT AVG(data)
FROM tbl
GROUP BY city
```

City	Count
NYC	7
Berkeley	5




ID	City	Data
1	NYC	0.78
2	NYC	0.13
3	Berkeley	0.25
4	NYC	0.19
5	NYC	0.11
6	Berkeley	0.09
7	NYC	0.18
8	NYC	0.15
9	Berkeley	0.13
10	Berkeley	0.49
11	NYC	0.19
12	Berkeley	0.10

Stratified samples

```
SELECT AVG(data)
FROM tbl
GROUP BY city
```

City	Count	Rate
NYC	7	2/7
Berkeley	5	2/5




ID	City	Data
1	NYC	0.78
2	NYC	0.13
3	Berkeley	0.25
4	NYC	0.19
5	NYC	0.11
6	Berkeley	0.09
7	NYC	0.18
8	NYC	0.15
9	Berkeley	0.13
10	Berkeley	0.49
11	NYC	0.19
12	Berkeley	0.10

Stratified samples

SELECT AVG(data)
FROM tbl
GROUP BY city


City	Count	Rate
NYC	7	2/7
Berkeley	5	2/5



ID	City	Data
1	NYC	0.78
2	NYC	0.13
3	Berkeley	0.25
4	NYC	0.19
5	NYC	0.11
6	Berkeley	0.09
7	NYC	0.18
8	NYC	0.15
9	Berkeley	0.13
10	Berkeley	0.49
11	NYC	0.19
12	Berkeley	0.10

Stratified samples

SELECT AVG(data)
FROM tbl
GROUP BY city



ID	City	Data	Rate
2	NYC	0.13	2/7
8	NYC	0.25	2/7
6	Berkeley	0.09	2/5
12	Berkeley	0.49	2/5

Does not change
query semantics

Are rare values represented in the sample now?

Sampling for rare QCS values

- When a QCS value is very rare, all records end up in the sample




Figure 4. Example of a stratified sample associated with a set of columns, ϕ .

What is BlinkDB?

- A framework that ...
 - creates and maintains a variety of uniform and stratified samples from underlying data
 - returns fast, approximate answers with error bars by executing queries on samples of data
 - verifies the correctness of the error bars that it returns at runtime

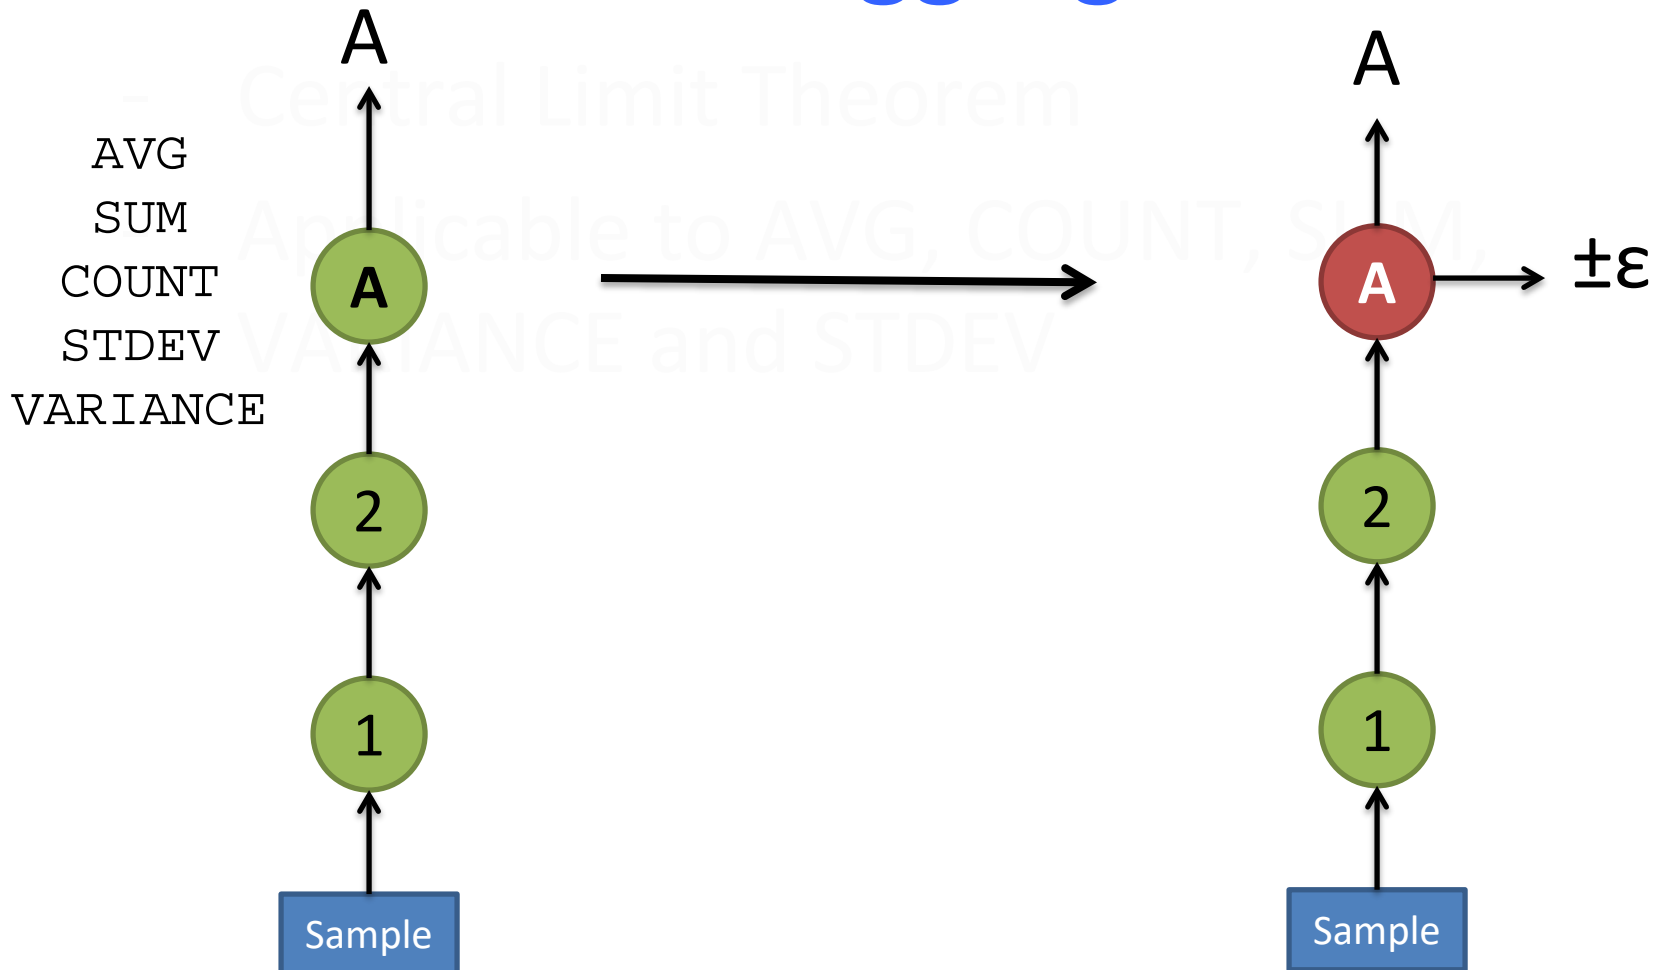
Error Estimation

- Closed Form Aggregate Functions

- Central Limit Theorem: When adding independent random variables, their sum tends towards a normal distribution, even if the original variables themselves are not normally distributed
- Applicable to AVG, COUNT, SUM, VARIANCE and STDEV

Error estimation

- Closed Form Aggregate Functions



Error estimation

- **Generalized** Aggregate Functions

- Statistical Bootstrap
- Applicable to complex and nested queries, UDFs, joins etc.

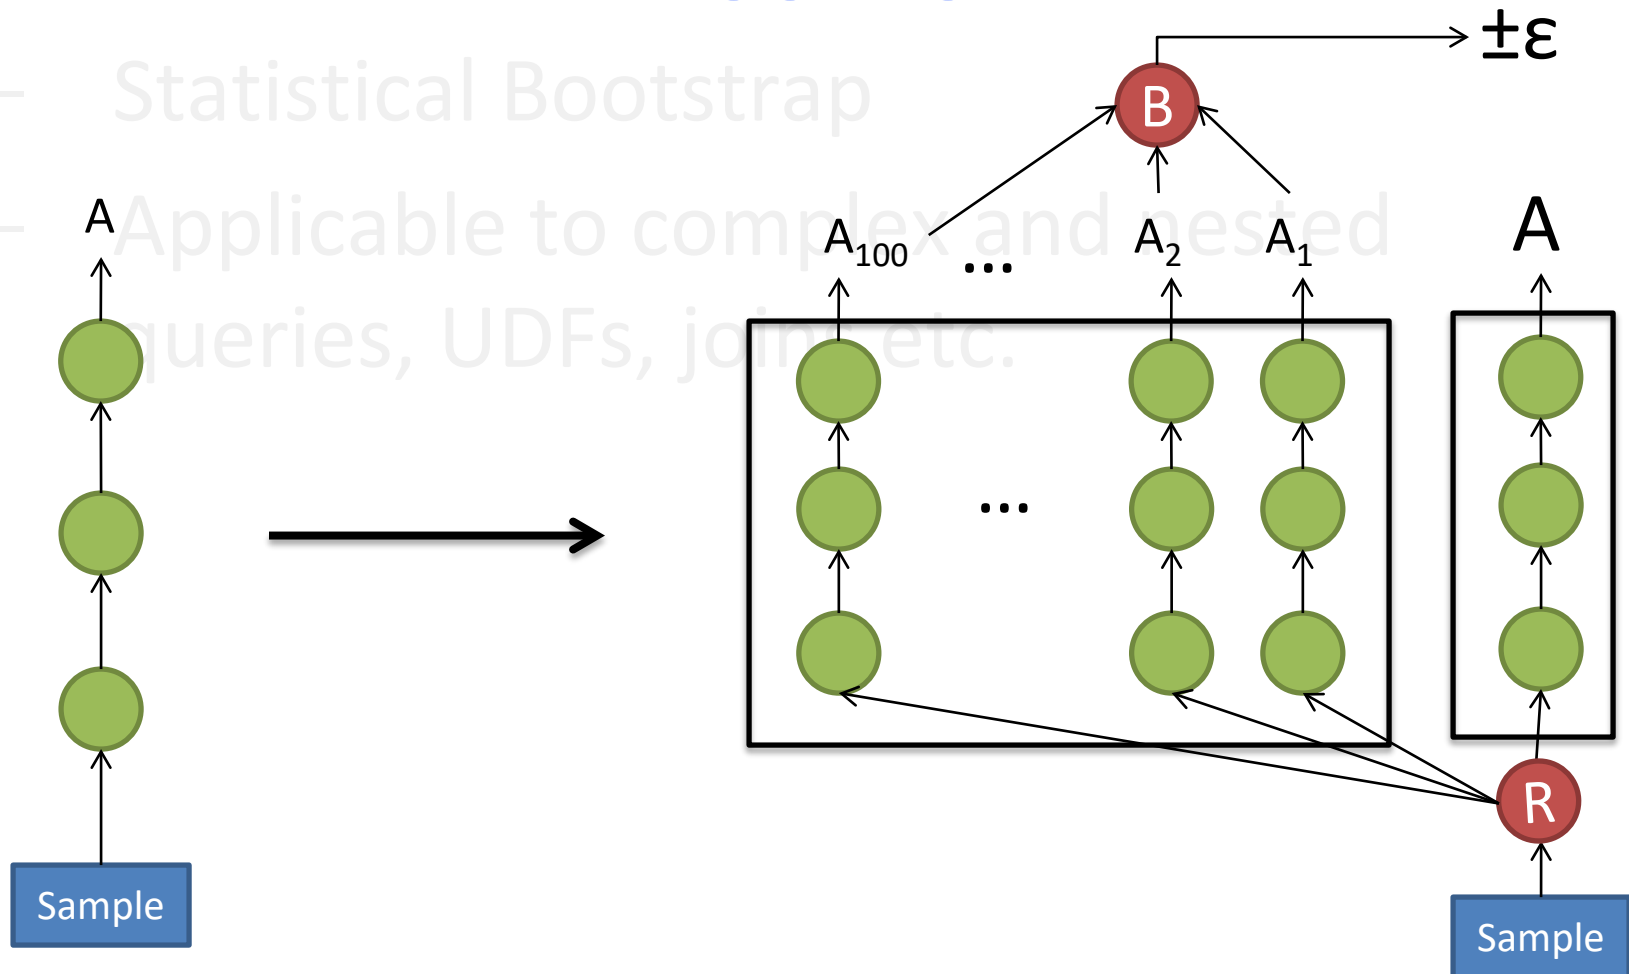
- Main idea

- Execute the query on x sub-samples
- Variance of answers: indicator of accuracy

Error estimation

• Generalized Aggregate Functions


- Statistical Bootstrap
- Applicable to complex and nested queries, UDFs, joins etc.



What is BlinkDB?

- A framework that ...
 - creates and maintains a variety of random and stratified samples from underlying data
 - returns fast, approximate answers with error bars by executing queries on samples of data
 - verifies the correctness of the error bars that it returns at runtime

Kleiner's Diagnostics



What is BlinkDB?

- A framework that ...
 - creates and maintains a variety of random and stratified samples from underlying data
 - returns fast, approximate answers with error bars by executing queries on samples of data
 - verifies the correctness of the error bars that it returns at runtime

Readings

- BlinkDB

- Sameer Agarwal, Barzan Mozafari, Aurojit Panda, Henry Milner, Samuel Madden, Ion Stoica: BlinkDB: queries with bounded errors and bounded response times on very large data. EuroSys 2013: 29-42
https://sameeragarwal.github.io/blinkdb_eurosys13.pdf
- Paper 2 (optional, will not be examined): <https://sameeragarwal.github.io/mod282-agarwal.pdf>

- Streams

- Matei Zaharia, Tathagata Das, Haoyuan Li, Timothy Hunter, Scott Shenker, Ion Stoica: Discretized streams: fault-tolerant streaming computation at scale. SOSP 2013: 423-438 http://people.csail.mit.edu/matei/papers/2013/sosp_spark_streaming.pdf
- Introduction & time models
 - Section 4 from <http://users.monash.edu/~mgaber/Muthu-Survey.pdf>
 - Section II.A from <http://dimacs.rutgers.edu/~graham/pubs/papers/fwddecay.pdf> for decay models
 - Section 3 from <http://dl.acm.org/citation.cfm?id=1060753> covers sliding windows, jumping windows, landmark windows

- Technical/How-to/API/dev...

- <http://spark.apache.org/sql/> <http://spark.apache.org/streaming/> <http://blinkdb.org>